

```
[1]: import pandas as pd
import numpy as np

from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report, roc_curve, auc, f1_score, precision_score, recall_score

import copy
import torch
import torch.nn as nn
from torch.optim import NAdam
import torch.nn.functional as F

import seaborn as sns
import matplotlib.pyplot as plt

from imblearn.over_sampling import RandomOverSampler

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")
```

Using device: cuda

```
[2]: diabetes=pd.read_csv('diabetes.csv')
```

```
[3]: diabetes.sample(5)
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
135	2	125	60	20	140	33.8	
677	0	93	60	0	0	35.3	
568	4	154	72	29	126	31.3	
356	1	125	50	40	167	33.3	
360	5	189	64	33	325	31.2	

	DiabetesPedigreeFunction	Age	Outcome
135	0.088	31	0
677	0.263	25	0
568	0.338	37	0
356	0.962	28	1
360	0.583	29	1

```
[4]: diabetes.nunique()
```

```
[4]: Pregnancies          17
      Glucose             136
```

```

BloodPressure      47
SkinThickness      51
Insulin            186
BMI                248
DiabetesPedigreeFunction  517
Age                52
Outcome            2
dtype: int64

```

```
[5]: diabetes.isnull().sum()
```

```

[5]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64

```

```

[6]: diabetes_x = diabetes.drop("Outcome", axis=1)
      diabetes_y = diabetes["Outcome"]

```

```
[7]: diabetes_x.describe()
```

```

[7]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count      768.000000  768.000000      768.000000      768.000000  768.000000
mean         3.845052  120.894531        69.105469        20.536458   79.799479
std          3.369578   31.972618        19.355807        15.952218  115.244002
min           0.000000    0.000000         0.000000         0.000000    0.000000
25%           1.000000   99.000000        62.000000         0.000000    0.000000
50%           3.000000  117.000000        72.000000        23.000000   30.500000
75%           6.000000  140.250000        80.000000        32.000000  127.250000
max          17.000000  199.000000       122.000000        99.000000  846.000000

      BMI  DiabetesPedigreeFunction      Age
count      768.000000          768.000000  768.000000
mean        31.992578           0.471876   33.240885
std          7.884160           0.331329   11.760232
min           0.000000           0.078000   21.000000
25%          27.300000           0.243750   24.000000
50%          32.000000           0.372500   29.000000
75%          36.600000           0.626250   41.000000
max          67.100000           2.420000   81.000000


```

```
[8]: sns.countplot(x='Outcome', data=diabetes, palette='Set2')
plt.title("Distribution of Diabetes Outcome")
plt.show()
```

/tmp/ipykernel_4973/2807747172.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Outcome', data=diabetes, palette='Set2')
```




```
[9]: diabetes_x_copy=diabetes_x.copy()
diabetes_y_copy=diabetes_y.copy()
```

```
[10]: rand_sampler = RandomOverSampler(sampling_strategy=1.0, random_state=42)
```

```
[11]: diabetes_x_resampled, diabetes_y_resampled = rand_sampler.
↳fit_resample(diabetes_x_copy, diabetes_y_copy)
```

```
[12]: plt.figure(figsize=(9, 5))
plt.pie(diabetes_y_resampled.value_counts(), labels=["No Diabetes", "Diabetes"],
        autopct="%1.1f%%", startangle=140)
plt.title("Distribution of Diabetes Cases After Oversampling")
plt.show()
```

Distribution of Diabetes Cases After Oversampling



```
[13]: print("Balanced dataset:")
print(f"Positive cases:{diabetes_x_resampled[diabetes_y_resampled == 1].
        shape[0]}\n Negative cases:{diabetes_x_resampled[diabetes_y_resampled == 0].
        shape[0]}")
```

```
Balanced dataset:
Positive cases:500
Negative cases:500
```

```
[14]: diabetes_x_resampled
```

```
[14]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1









4	0	137	40	35	168	43.1
..
995	1	122	64	32	156	35.1
996	0	131	0	0	0	43.2
997	8	120	0	0	0	30.0
998	4	111	72	47	207	37.1
999	13	158	114	0	0	42.3

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..
995	0.692	30
996	0.270	26
997	0.183	38
998	1.390	56
999	0.257	44

[1000 rows x 8 columns]

```
[15]: diabetes_combined=pd.concat([diabetes_x_resampled,diabetes_y_resampled],axis=1)
```

```
[32]: fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(11, 20))
for i, col in enumerate(diabetes_combined.columns[:-1]):
    ax = axes[i // 2, i % 2]
    sns.boxplot(data=diabetes_combined, x="Outcome", y=col, ax=ax)
```



```
[20]: def remove_outliers(df, columns):
        for col in columns:
            Q1 = df[col].quantile(0.25)
            Q3 = df[col].quantile(0.75)
            IQR = Q3 - Q1
            lower_bound = Q1 - 1.5 * IQR
            upper_bound = Q3 + 1.5 * IQR
            df[(df[col] < lower_bound)] = pd.NA
            df[(df[col] > upper_bound)] = pd.NA
        return df

[21]: diabetes_x_resampled = remove_outliers(diabetes_x_resampled,
        ↪diabetes_x_resampled.columns)









[22]: diabetes_x_copy = remove_outliers(diabetes_x_copy, diabetes_x_copy.columns)

[23]: combined_df = pd.concat([diabetes_x_resampled, diabetes_y_resampled], axis=1)
        combined_df_cp = pd.concat([diabetes_x_copy, diabetes_y_copy], axis=1)

[24]: combined_df.dropna(inplace=True)
        combined_df_cp.dropna(inplace=True)

[25]: diabetes_x = combined_df.drop("Outcome", axis=1)
        diabetes_y = combined_df["Outcome"]
        diabetes_x_copy = combined_df_cp.drop("Outcome", axis=1)
        diabetes_y_copy = combined_df_cp["Outcome"]

[33]: fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(11, 20))
        for i, col in enumerate(diabetes_x.columns):
            ax = axes[i // 2, i % 2]
            sns.boxplot(data=diabetes_x, x=diabetes_y, y=col, ax=ax)
```





```
[34]: print(f"After removing outliers, there are {diabetes_x.shape[0]} samples left.")
      print(f"There are {diabetes_x[diabetes_y == 1].shape[0]} positive cases and_
      ↪{diabetes_x[diabetes_y== 0].shape[0]} negative cases")
```

After removing outliers, there are 835 samples left.
There are 394 positive cases and 441 negative cases

Correlation matrix of balanced dataset

```
[43]: corr = combined_df.corr()
      plt.figure(figsize=(9, 9))
      sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm", square=True)
```


[43]: <Axes: >



Correlation matrix of unbalanced dataset

```
[45]: corr_cp = combined_df_cp.corr()
plt.figure(figsize=(9, 9))
sns.heatmap(corr_cp, annot=True, fmt=".2f", cmap="coolwarm", square=True)
```

[45]: <Axes: >



```
[46]: scaler = StandardScaler()
diabetes_X_scaled = scaler.fit_transform(diabetes_x)
diabetes_X_scaled_cp = scaler.fit_transform(diabetes_x_copy)
```

```
[47]: X_train, X_test, y_train, y_test = train_test_split(diabetes_X_scaled, ↵  
↪diabetes_y, test_size=0.2, random_state=42)
```

```
[48]: X_train_02, X_test_02, y_train_02, y_test_02 = ↵  
↪train_test_split(diabetes_X_scaled_cp, diabetes_y_copy, test_size=0.2, ↵  
↪random_state=42)
```

Training a Random Forest Classifier on the balanced dataset

```
[50]: rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
rf_model.fit(X_train, y_train)
```

```
[50]: RandomForestClassifier(random_state=42)
```

```
[51]: y_pred = rf_model.predict(X_test)
```

```
[52]: accuracy = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy)
```


Accuracy: 0.8802395209580839

```
[53]: confusion_mat = confusion_matrix(y_test, y_pred)  
print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[75 13]  
 [ 7 72]]
```

```
[56]: plt.figure(figsize=(6, 6))  
sns.heatmap(confusion_mat, annot=True, annot_kws={"size": 16}, fmt='d', ↵  
↪cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No ↵  
↪Diabetes", "Diabetes"], square=True)  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.title("Confusion Matrix (Random Forest)")  
plt.show()
```



```
[57]: print("Classification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.85	0.88	88
1	0.85	0.91	0.88	79
accuracy			0.88	167
macro avg	0.88	0.88	0.88	167
weighted avg	0.88	0.88	0.88	167

Training a Random Forest Classifier on the unbalanced dataset

```
[59]: rf_model_02 = RandomForestClassifier(n_estimators=100, random_state=42)
      rf_model_02.fit(X_train_02, y_train_02)
```

```
[59]: RandomForestClassifier(random_state=42)
```

```
[60]: y_pred_02 = rf_model_02.predict(X_test_02)
```

```
[61]: accuracy_02 = accuracy_score(y_test_02, y_pred_02)
print("Accuracy on unbalanced dataset:", accuracy_02)
```

Accuracy on unbalanced dataset: 0.7421875

```
[62]: y_test_02
```


```
[62]: 338    1
      763    0
      104    0
      437    0
      184    0
      ..
      723    0
      246    0
      346    0
      275    0
      358    0
      Name: Outcome, Length: 128, dtype: int64
```

```
[63]: confusion_matrix_02 = confusion_matrix(y_test_02, y_pred_02)
print("Confusion Matrix (Unbalanced Dataset):")
print(confusion_matrix_02)
```

Confusion Matrix (Unbalanced Dataset):

```
[[79 13]
 [20 16]]
```

```
[64]: plt.figure(figsize=(6,6))
sns.heatmap(confusion_matrix_02, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No
            Diabetes", "Diabetes"], square=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix (Random Forest)")
plt.show()
```



```
[65]: print("Classification Report:\n", classification_report(y_test_02, y_pred_02))
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.86	0.83	92
1	0.55	0.44	0.49	36
accuracy			0.74	128
macro avg	0.67	0.65	0.66	128
weighted avg	0.73	0.74	0.73	128

Training a Multi-Layer Perceptron on the balanced dataset

```
[142]: class DiabetesModel(nn.Module):
        def __init__(self, input_size=8, hidden_size=14, output_size=1):
            super(DiabetesModel, self).__init__()
            self.fc1 = nn.Linear(input_size, hidden_size)
```

```

        self.fc2 = nn.Linear(hidden_size, hidden_size)
        self.out = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        x = F.leaky_relu(self.fc1(x))
        x = F.leaky_relu(self.fc2(x))
        x = self.out(x)
        return x

```

```

[143]: torch.manual_seed(42)
nn_model_01 = DiabetesModel()
nn_model_01
nn_model_01.to(device=device)

```

```

[143]: DiabetesModel(
  (fc1): Linear(in_features=8, out_features=14, bias=True)
  (fc2): Linear(in_features=14, out_features=14, bias=True)
  (out): Linear(in_features=14, out_features=1, bias=True)
)

```

```

[144]: torch.manual_seed(42)
nn_model_02 = DiabetesModel()
nn_model_02
nn_model_02.to(device=device)

```

```

[144]: DiabetesModel(
  (fc1): Linear(in_features=8, out_features=14, bias=True)
  (fc2): Linear(in_features=14, out_features=14, bias=True)
  (out): Linear(in_features=14, out_features=1, bias=True)
)

```

We are using cross entropy loss for calculating the loss of the model, and NADAM for optimizing the weights and biases of the model

```

[145]: loss_fn = nn.BCEWithLogitsLoss()
optimizer = NAdam(nn_model_01.parameters(), lr=0.01)

```

```

[146]: loss_fn2 = nn.BCEWithLogitsLoss()
optimizer2 = NAdam(nn_model_02.parameters(), lr=0.01)

```

```

[147]: X_train_tensor = torch.tensor(X_train, dtype=torch.float32, device=device)
y_train_tensor = torch.tensor(y_train.values, dtype=torch.float32, ↵
↵device=device).view(-1, 1)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32, device=device)
y_test_tensor = torch.tensor(y_test.values, dtype=torch.float32, device=device).
↵view(-1, 1)

```

```
[148]: X_train_tensor1 = torch.tensor(X_train_02, dtype=torch.float32, device=device)
y_train_tensor1 = torch.tensor(y_train_02.values, dtype=torch.float32,
    ↪device=device).view(-1, 1)
X_test_tensor1 = torch.tensor(X_test_02, dtype=torch.float32, device=device)
y_test_tensor1 = torch.tensor(y_test_02.values, dtype=torch.float32,
    ↪device=device).view(-1, 1)
```

Training the neural network on the balanced dataset

```
[149]: epochs = 1000
best_model_loss = float('inf')
best_model_weights = None
patience = 7

loss_list = []
accuracy_list = []

for i in range(epochs):
    # train_loss = 0.0
    nn_model_01.train()

    outputs = nn_model_01(X_train_tensor)
    loss = loss_fn(outputs, y_train_tensor)
    loss_list.append(loss.item())

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    nn_model_01.eval()
    with torch.no_grad():
        outputs = nn_model_01(X_test_tensor)
        test_loss = loss_fn(outputs, y_test_tensor)
    y_pred = torch.sigmoid(outputs).round()
    accuracy = (y_pred == y_test_tensor).float().mean().item()
    accuracy_list.append(accuracy)

    if test_loss < best_model_loss:
        best_model_loss = test_loss
        best_model_weights = copy.deepcopy(nn_model_01.state_dict())
        patience_counter = 0
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print(f"Early stopping at epoch {i+1}")
            break
```



```
print(f"Epoch {i+1}/{epochs}, Train Loss: {loss.item():.4f}, Val Loss:␣  
↪{test_loss.item():.4f}")
```

```
Epoch 1/1000, Train Loss: 0.6923, Val Loss: 0.6869  
Epoch 2/1000, Train Loss: 0.6869, Val Loss: 0.6823  
Epoch 3/1000, Train Loss: 0.6820, Val Loss: 0.6764  
Epoch 4/1000, Train Loss: 0.6764, Val Loss: 0.6688  
Epoch 5/1000, Train Loss: 0.6694, Val Loss: 0.6588  
Epoch 6/1000, Train Loss: 0.6604, Val Loss: 0.6463  
Epoch 7/1000, Train Loss: 0.6490, Val Loss: 0.6315  
Epoch 8/1000, Train Loss: 0.6355, Val Loss: 0.6143  
Epoch 9/1000, Train Loss: 0.6199, Val Loss: 0.5949  
Epoch 10/1000, Train Loss: 0.6030, Val Loss: 0.5750  
Epoch 11/1000, Train Loss: 0.5853, Val Loss: 0.5558  
Epoch 12/1000, Train Loss: 0.5682, Val Loss: 0.5389  
Epoch 13/1000, Train Loss: 0.5522, Val Loss: 0.5242  
Epoch 14/1000, Train Loss: 0.5377, Val Loss: 0.5100  
Epoch 15/1000, Train Loss: 0.5244, Val Loss: 0.4983  
Epoch 16/1000, Train Loss: 0.5125, Val Loss: 0.4879  
Epoch 17/1000, Train Loss: 0.5018, Val Loss: 0.4791  
Epoch 18/1000, Train Loss: 0.4919, Val Loss: 0.4711  
Epoch 19/1000, Train Loss: 0.4831, Val Loss: 0.4650  
Epoch 20/1000, Train Loss: 0.4754, Val Loss: 0.4592  
Epoch 21/1000, Train Loss: 0.4683, Val Loss: 0.4551  
Epoch 22/1000, Train Loss: 0.4621, Val Loss: 0.4507  
Epoch 23/1000, Train Loss: 0.4567, Val Loss: 0.4507  
Epoch 24/1000, Train Loss: 0.4522, Val Loss: 0.4470  
Epoch 25/1000, Train Loss: 0.4484, Val Loss: 0.4493  
Epoch 26/1000, Train Loss: 0.4449, Val Loss: 0.4447  
Epoch 27/1000, Train Loss: 0.4420, Val Loss: 0.4546  
Epoch 28/1000, Train Loss: 0.4407, Val Loss: 0.4447  
Epoch 29/1000, Train Loss: 0.4410, Val Loss: 0.4588  
Epoch 30/1000, Train Loss: 0.4378, Val Loss: 0.4443  
Epoch 31/1000, Train Loss: 0.4347, Val Loss: 0.4543  
Epoch 32/1000, Train Loss: 0.4310, Val Loss: 0.4452  
Epoch 33/1000, Train Loss: 0.4284, Val Loss: 0.4515  
Epoch 34/1000, Train Loss: 0.4260, Val Loss: 0.4448  
Epoch 35/1000, Train Loss: 0.4240, Val Loss: 0.4508  
Epoch 36/1000, Train Loss: 0.4219, Val Loss: 0.4430  
Epoch 37/1000, Train Loss: 0.4201, Val Loss: 0.4496  
Epoch 38/1000, Train Loss: 0.4185, Val Loss: 0.4418  
Epoch 39/1000, Train Loss: 0.4170, Val Loss: 0.4500  
Epoch 40/1000, Train Loss: 0.4154, Val Loss: 0.4406  
Epoch 41/1000, Train Loss: 0.4141, Val Loss: 0.4509  
Epoch 42/1000, Train Loss: 0.4127, Val Loss: 0.4395  
Epoch 43/1000, Train Loss: 0.4117, Val Loss: 0.4519  
Epoch 44/1000, Train Loss: 0.4096, Val Loss: 0.4389  
Epoch 45/1000, Train Loss: 0.4085, Val Loss: 0.4501
```

```
Epoch 46/1000, Train Loss: 0.4065, Val Loss: 0.4391
Epoch 47/1000, Train Loss: 0.4052, Val Loss: 0.4487
Epoch 48/1000, Train Loss: 0.4036, Val Loss: 0.4384
Epoch 49/1000, Train Loss: 0.4028, Val Loss: 0.4481
Epoch 50/1000, Train Loss: 0.4012, Val Loss: 0.4380
Epoch 51/1000, Train Loss: 0.4000, Val Loss: 0.4475
Epoch 52/1000, Train Loss: 0.3981, Val Loss: 0.4370
Epoch 53/1000, Train Loss: 0.3964, Val Loss: 0.4458
Epoch 54/1000, Train Loss: 0.3949, Val Loss: 0.4355
Epoch 55/1000, Train Loss: 0.3940, Val Loss: 0.4448
Epoch 56/1000, Train Loss: 0.3924, Val Loss: 0.4342
Epoch 57/1000, Train Loss: 0.3914, Val Loss: 0.4448
Epoch 58/1000, Train Loss: 0.3905, Val Loss: 0.4322
Epoch 59/1000, Train Loss: 0.3901, Val Loss: 0.4444
Epoch 60/1000, Train Loss: 0.3891, Val Loss: 0.4324
Epoch 61/1000, Train Loss: 0.3894, Val Loss: 0.4460
Epoch 62/1000, Train Loss: 0.3875, Val Loss: 0.4316
Epoch 63/1000, Train Loss: 0.3868, Val Loss: 0.4446
Epoch 64/1000, Train Loss: 0.3850, Val Loss: 0.4306
Epoch 65/1000, Train Loss: 0.3848, Val Loss: 0.4442
Epoch 66/1000, Train Loss: 0.3825, Val Loss: 0.4299
Epoch 67/1000, Train Loss: 0.3805, Val Loss: 0.4418
Epoch 68/1000, Train Loss: 0.3794, Val Loss: 0.4301
Epoch 69/1000, Train Loss: 0.3789, Val Loss: 0.4418
Epoch 70/1000, Train Loss: 0.3769, Val Loss: 0.4303
Epoch 71/1000, Train Loss: 0.3762, Val Loss: 0.4434
Epoch 72/1000, Train Loss: 0.3755, Val Loss: 0.4307
Early stopping at epoch 73
```

```
[150]: nn_model_01.to(device="cpu")
nn_model_01.load_state_dict(best_model_weights)
nn_model_01.eval()

X_test_tensor = X_test_tensor.to(device="cpu")

with torch.no_grad():
    outputs = nn_model_01(X_test_tensor)
    y_pred = torch.sigmoid(outputs).round().numpy()
```

```
[151]: y_test_tensor = y_test_tensor.to(device="cpu")

print("Classification Report:")
print(classification_report(y_test_tensor, y_pred))
```


```
Classification Report:
              precision    recall  f1-score   support

0.0           0.00         0.84         0.85         88
```

	1.0	0.83	0.85	0.84	79
accuracy				0.84	167
macro avg	0.84	0.84	0.84	0.84	167
weighted avg	0.84	0.84	0.84	0.84	167

```
[152]: conf_matrix = confusion_matrix(y_test_tensor, y_pred)
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=['No Diabetes', 'Diabetes'], yticklabels=['No
            Diabetes', 'Diabetes'], square=True)
```


[152]: <Axes: >



```
[153]: print(f"Accuracy: {accuracy_score(y_test_tensor, y_pred):.4f}")
```

Accuracy: 0.8443

```
[156]: plt.figure(figsize=(10, 5))
sns.lineplot(x=range(len(loss_list)), y=loss_list, label='Train Loss')
sns.lineplot(x=range(len(accuracy_list)), y=accuracy_list, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss / Accuracy')
plt.title('Loss and Accuracy over Epochs')
plt.legend()
plt.show()
```



Training the neural network on the unbalanced dataset

```
[157]: epochs = 400
best_model_loss1 = float('inf')
best_model_weights1 = None
patience = 7

loss_list1 = []
accuracy_list1 = []

for i in range(epochs):
    # train_loss = 0.0
    nn_model_02.train()

    outputs1 = nn_model_02(X_train_tensor1)
    loss1 = loss_fn2(outputs1, y_train_tensor1)
    loss_list1.append(loss1.item())
```

```

optimizer2.zero_grad()
loss1.backward()
optimizer2.step()

nn_model_02.eval()
with torch.no_grad():
    outputs1 = nn_model_02(X_test_tensor1)
    test_loss1 = loss_fn2(outputs1, y_test_tensor1)
y_pred = torch.sigmoid(outputs1).round()
accuracy = (y_pred == y_test_tensor1).float().mean().item()
accuracy_list1.append(accuracy)

if test_loss1 < best_model_loss1:
    best_model_loss1 = test_loss1
    best_model_weights1 = copy.deepcopy(nn_model_02.state_dict())
    patience_counter1 = 0
else:
    patience_counter1 += 1
    if patience_counter1 >= patience:
        print(f"Early stopping at epoch {i+1}")
        break
print(f"Epoch {i+1}/{epochs}, Train Loss: {loss1.item():.4f}, Val Loss: ↵
↵{test_loss1.item():.4f}")

```

```

Epoch 1/400, Train Loss: 0.6982, Val Loss: 0.6940
Epoch 2/400, Train Loss: 0.6881, Val Loss: 0.6849
Epoch 3/400, Train Loss: 0.6794, Val Loss: 0.6753
Epoch 4/400, Train Loss: 0.6702, Val Loss: 0.6646
Epoch 5/400, Train Loss: 0.6596, Val Loss: 0.6520
Epoch 6/400, Train Loss: 0.6469, Val Loss: 0.6371
Epoch 7/400, Train Loss: 0.6317, Val Loss: 0.6201
Epoch 8/400, Train Loss: 0.6138, Val Loss: 0.6022
Epoch 9/400, Train Loss: 0.5938, Val Loss: 0.5833
Epoch 10/400, Train Loss: 0.5725, Val Loss: 0.5659
Epoch 11/400, Train Loss: 0.5512, Val Loss: 0.5519
Epoch 12/400, Train Loss: 0.5309, Val Loss: 0.5404
Epoch 13/400, Train Loss: 0.5126, Val Loss: 0.5316
Epoch 14/400, Train Loss: 0.4969, Val Loss: 0.5252
Epoch 15/400, Train Loss: 0.4829, Val Loss: 0.5201
Epoch 16/400, Train Loss: 0.4703, Val Loss: 0.5167
Epoch 17/400, Train Loss: 0.4592, Val Loss: 0.5143
Epoch 18/400, Train Loss: 0.4492, Val Loss: 0.5122
Epoch 19/400, Train Loss: 0.4401, Val Loss: 0.5112
Epoch 20/400, Train Loss: 0.4317, Val Loss: 0.5104
Epoch 21/400, Train Loss: 0.4238, Val Loss: 0.5094
Epoch 22/400, Train Loss: 0.4161, Val Loss: 0.5070
Epoch 23/400, Train Loss: 0.4087, Val Loss: 0.5066
Epoch 24/400, Train Loss: 0.4025, Val Loss: 0.5068

```

```
Epoch 25/400, Train Loss: 0.3970, Val Loss: 0.5074
Epoch 26/400, Train Loss: 0.3921, Val Loss: 0.5081
Epoch 27/400, Train Loss: 0.3879, Val Loss: 0.5087
Epoch 28/400, Train Loss: 0.3842, Val Loss: 0.5107
Epoch 29/400, Train Loss: 0.3809, Val Loss: 0.5101
Early stopping at epoch 30
```

```
[158]: nn_model_02.load_state_dict(best_model_weights1)
nn_model_02.to(device="cpu")
X_test_tensor1 = X_test_tensor1.to(device="cpu")

nn_model_02.eval()
with torch.no_grad():
    outputs1 = nn_model_02(X_test_tensor1)
    y_pred_02 = torch.sigmoid(outputs1).round().numpy()
```

```
[159]: y_test_tensor1 = y_test_tensor1.to(device="cpu")

print("Classification Report:")
print(classification_report(y_test_tensor1, y_pred_02))
```


Classification Report:

	precision	recall	f1-score	support
0.0	0.80	0.88	0.84	92
1.0	0.59	0.44	0.51	36
accuracy			0.76	128
macro avg	0.70	0.66	0.67	128
weighted avg	0.74	0.76	0.75	128

```
[160]: print(f"Accuracy: {accuracy_score(y_test_tensor1, y_pred_02):.4f}")
```

Accuracy: 0.7578

```
[161]: plt.figure(figsize=(10, 5))
sns.lineplot(x=range(len(loss_list1)), y=loss_list1, label='Train Loss')
sns.lineplot(x=range(len(accuracy_list1)), y=accuracy_list1, label='Test_
↳Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss / Accuracy')
plt.title('Loss and Accuracy over Epochs')
plt.legend()
plt.show()
```



Training a Support Vector Machine on the balanced dataset

```
[162]: svc_model = SVC(random_state=42, kernel='rbf')
```

```
[163]: svc_model.fit(X_train, y_train)
```

```
[163]: SVC(random_state=42)
```

```
[164]: y_pred = svc_model.predict(X_test)
```

```
[165]: accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
```

Accuracy: 0.8084


```
[166]: confusion_mat = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[73 15]
 [17 62]]
```

```
[167]: plt.figure(figsize=(6, 6))
sns.heatmap(confusion_mat, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No
            Diabetes", "Diabetes"], square=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
plt.title("Confusion Matrix (SVM)")
plt.show()
```



```
[168]: print("Classification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.83	0.82	88
1	0.81	0.78	0.79	79
accuracy			0.81	167
macro avg	0.81	0.81	0.81	167
weighted avg	0.81	0.81	0.81	167

Training a Support Vector Machine on the unbalanced dataset.

```
[169]: svc_model_02 = SVC(random_state=42, kernel='rbf')
```



```
[170]: svc_model_02.fit(X_train_02, y_train_02)
```

```
[170]: SVC(random_state=42)
```

```
[171]: y_pred_02 = svc_model_02.predict(X_test_02)
```

```
[172]: accuracy_02 = accuracy_score(y_test_02, y_pred_02)
print(f"Accuracy: {accuracy_02:.4f}")
```


Accuracy: 0.7578

```
[173]: confusion_mat = confusion_matrix(y_test_02, y_pred_02)
print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[81 11]
 [20 16]]
```

```
[174]: plt.figure(figsize=(6, 6))
sns.heatmap(confusion_mat, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No
            Diabetes", "Diabetes"], square=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix (SVM)")
plt.show()
```



```
[175]: print("Classification Report:\n", classification_report(y_test_02, y_pred_02))
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.88	0.84	92
1	0.59	0.44	0.51	36
accuracy			0.76	128
macro avg	0.70	0.66	0.67	128
weighted avg	0.74	0.76	0.75	128

Plotting the ROC-AUC curves of the models trained on the balanced dataset

```
[176]: y_pred_rf = rf_model.predict_proba(X_test)[: , 1]
y_pred_svc = svc_model.decision_function(X_test)
y_pred_nn = torch.sigmoid(nn_model_01(X_test_tensor)).detach().numpy()
```

```


[179]: plt.figure(figsize=(7, 7))

fpr1, tpr1, _ = roc_curve(y_test, y_pred_rf)
fpr2, tpr2, _ = roc_curve(y_test, y_pred_svc)
fpr3, tpr3, _ = roc_curve(y_test, y_pred_nn)

roc_auc_rf = auc(fpr1, tpr1)
roc_auc_svc = auc(fpr2, tpr2)
roc_auc_nn = auc(fpr3, tpr3)

plt.plot(fpr1, tpr1, color='blue', label=f'Random Forest (AUC = {roc_auc_rf:.2f})')
plt.plot(fpr2, tpr2, color='green', label=f'SVM (AUC = {roc_auc_svc:.2f})')
plt.plot(fpr3, tpr3, color='red', label=f'Neural Network (AUC = {roc_auc_nn:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

```



```
[180]: nn_model_01.eval()
with torch.no_grad():
    outputs = nn_model_01(X_test_tensor)
    y_pred = torch.sigmoid(outputs).round().numpy()

results_df = pd.DataFrame({
    'Model': ['Random Forest', 'Neural Network', 'SVM'],
    'Accuracy': [
        accuracy_score(y_test, rf_model.predict(X_test)),
        accuracy_score(y_test, y_pred),
        accuracy_score(y_test, svc_model.predict(X_test)),
    ],
    'AUC': [roc_auc_rf, roc_auc_svc, roc_auc_nn],
```

```

    'F1-Score': [
        f1_score(y_test, rf_model.predict(X_test)),
        f1_score(y_test, y_pred),
        f1_score(y_test, svc_model.predict(X_test)),
    ],
    'Precision': [
        precision_score(y_test, rf_model.predict(X_test)),
        precision_score(y_test, y_pred),
        precision_score(y_test, svc_model.predict(X_test)),
    ],
    'Recall': [
        recall_score(y_test, rf_model.predict(X_test)),
        recall_score(y_test, y_pred),
        recall_score(y_test, svc_model.predict(X_test)),
    ]
})

```

```
[181]: results_df
```

```
[181]:
```

	Model	Accuracy	AUC	F1-Score	Precision	Recall
0	Random Forest	0.880240	0.946778	0.878049	0.847059	0.911392
1	Neural Network	0.844311	0.889097	0.837500	0.827160	0.848101
2	SVM	0.808383	0.886939	0.794872	0.805195	0.784810