

# Pattern Recognition Laboratory

Mini Project  
Report

Submitted By

Aswin Tom  
(97424607006)

Hrishikesh M Jayakrishnan  
(97424607010)

Mohammad Farhan Z  
(97424607012)

MSc Computer Science with Specialization in Artificial Intelligence

## Mini Project

### Aim:

To train a neural network classifier using the PIMA Indian Diabetes dataset by performing appropriate data preprocessing, model training, and performance evaluation, and to compare its results with at least two traditional machine learning models using metrics such as accuracy and confusion matrix.

### Program Code:

```
[1]: import pandas as pd
import numpy as np

from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, \
    →classification_report, roc_curve, auc, f1_score, precision_score, \
    →recall_score

import copy
import torch
import torch.nn as nn
from torch.optim import NAdam
import torch.nn.functional as F

import seaborn as sns
import matplotlib.pyplot as plt

from imblearn.over_sampling import RandomOverSampler

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")
```

Using device: cuda

```
[2]: diabetes=pd.read_csv('diabetes.csv')
```

```
[3]: diabetes.sample(5)
```

```
[3]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
135             2     125             60             20     140  33.8
677             0      93             60              0         0  35.3
```

568	4	154	72	29	126	31.3
356	1	125	50	40	167	33.3
360	5	189	64	33	325	31.2

	DiabetesPedigreeFunction	Age	Outcome
135	0.088	31	0
677	0.263	25	0
568	0.338	37	0
356	0.962	28	1
360	0.583	29	1

```
[4]: diabetes.nunique()
```

```
[4]: Pregnancies      17
      Glucose          136
      BloodPressure    47
      SkinThickness    51
      Insulin          186
      BMI              248
      DiabetesPedigreeFunction  517
      Age              52
      Outcome           2
      dtype: int64
```

```
[5]: diabetes.isnull().sum()
```

```
[5]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome           0
      dtype: int64
```

```
[6]: diabetes_x = diabetes.drop("Outcome", axis=1)
      diabetes_y = diabetes["Outcome"]
```

```
[7]: diabetes_x.describe()
```

```
[7]:      Pregnancies      Glucose  BloodPressure  SkinThickness  Insulin  \
count      768.000000  768.000000    768.000000    768.000000  768.000000
mean         3.845052  120.894531     69.105469     20.536458    79.799479
```

std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000


	BMI	DiabetesPedigreeFunction	Age
count	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885
std	7.884160	0.331329	11.760232
min	0.000000	0.078000	21.000000
25%	27.300000	0.243750	24.000000
50%	32.000000	0.372500	29.000000
75%	36.600000	0.626250	41.000000
max	67.100000	2.420000	81.000000

```
[8]: sns.countplot(x='Outcome', data=diabetes, palette='Set2')
plt.title("Distribution of Diabetes Outcome")
plt.show()
```

/tmp/ipykernel\_4973/2807747172.py:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Outcome', data=diabetes, palette='Set2')
```




```
[9]: diabetes_x_copy=diabetes_x.copy()  
     diabetes_y_copy=diabetes_y.copy()
```

```
[10]: rand_sampler = RandomOverSampler(sampling_strategy=1.0, random_state=42)
```

```
[11]: diabetes_x_resampled, diabetes_y_resampled = rand_sampler.  
     ↪ fit_resample(diabetes_x_copy, diabetes_y_copy)
```

```
[12]: plt.figure(figsize=(9, 5))  
     plt.pie(diabetes_y_resampled.value_counts(), labels=["No Diabetes",  
     ↪ "Diabetes"], autopct="%1.1f%%", startangle=140)  
     plt.title("Distribution of Diabetes Cases After Oversampling")  
     plt.show()
```

Distribution of Diabetes Cases After Oversampling



```
[13]: print("Balanced dataset:")
print(f"Positive cases:{diabetes_x_resampled[diabetes_y_resampled == 1].
      ↳shape[0]}\n Negative cases:{diabetes_x_resampled[diabetes_y_resampled_
      ↳== 0].shape[0]}")
```

```
Balanced dataset:
Positive cases:500
Negative cases:500
```

```
[14]: diabetes_x_resampled
```

```
[14]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
995	1	122	64	32	156	35.1	
996	0	131	0	0	0	43.2	

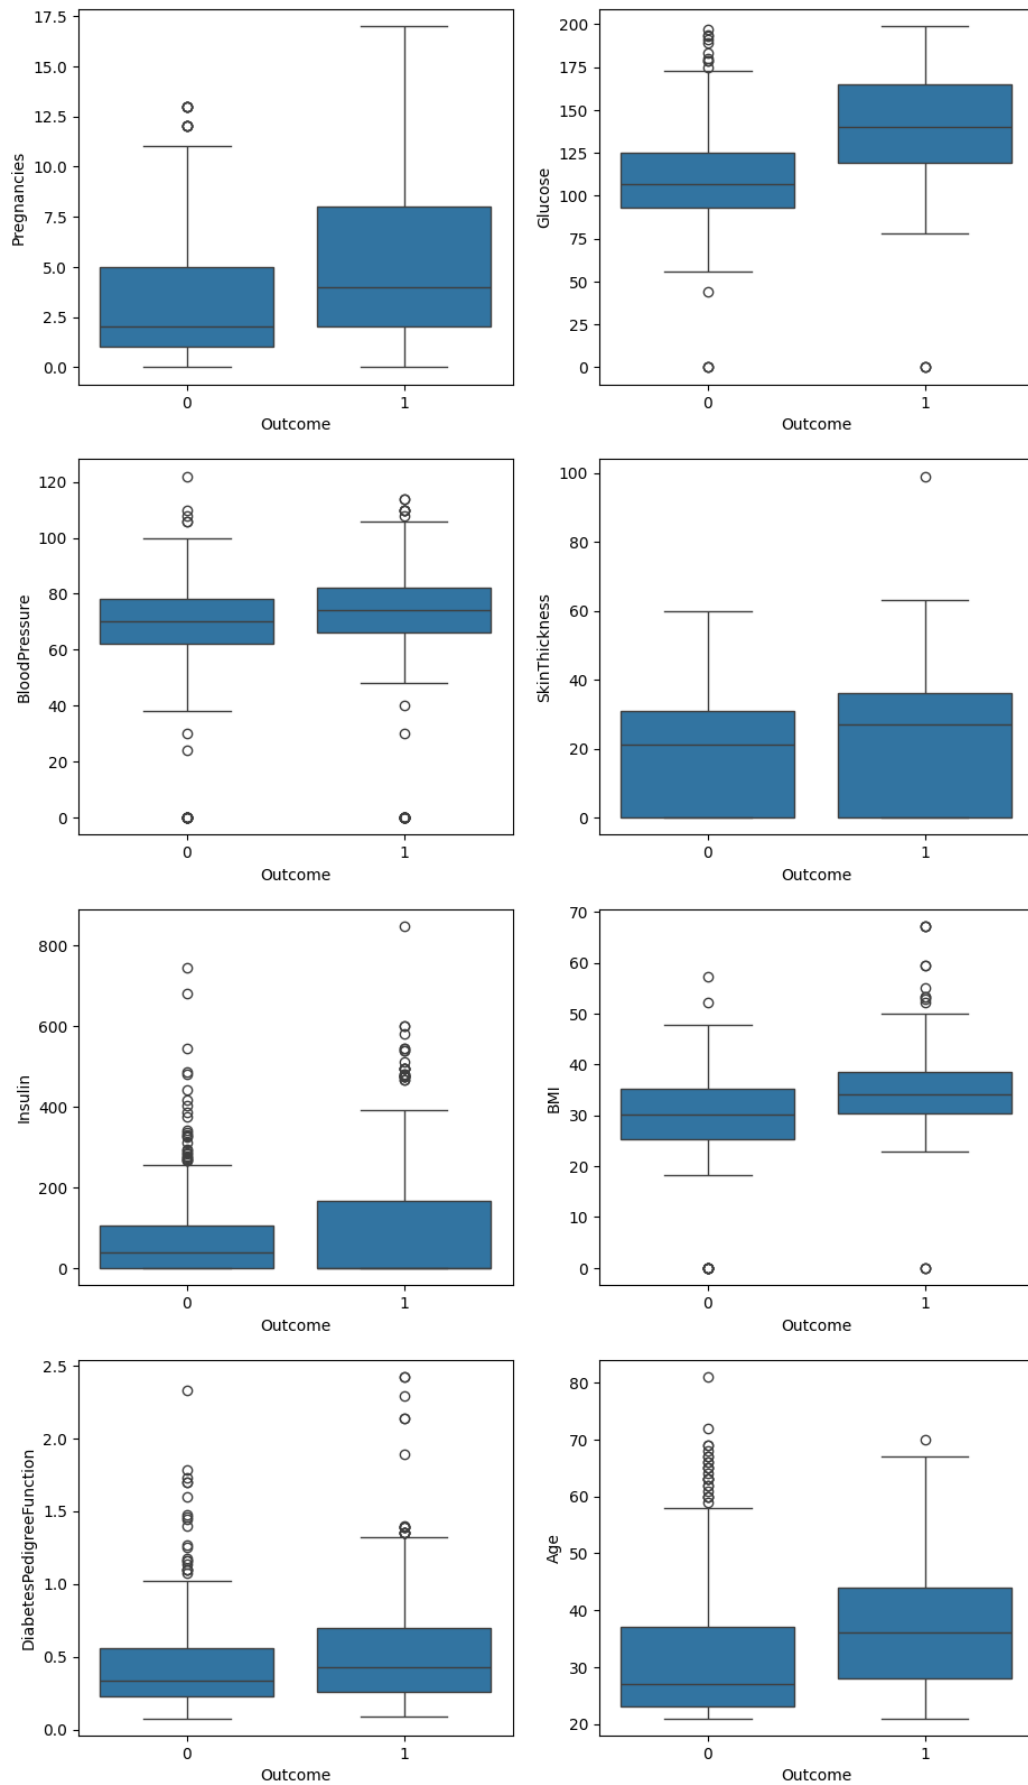
997	8	120	0	0	0	30.0
998	4	111	72	47	207	37.1
999	13	158	114	0	0	42.3

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..	...	...
995	0.692	30
996	0.270	26
997	0.183	38
998	1.390	56
999	0.257	44

[1000 rows x 8 columns]

```
[15]: diabetes_combined=pd.
      ↪concat([diabetes_x_resampled,diabetes_y_resampled],axis=1)
```

```
[32]: fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(11, 20))
      for i, col in enumerate(diabetes_combined.columns[:-1]):
          ax = axes[i // 2, i % 2]
          sns.boxplot(data=diabetes_combined, x="Outcome", y=col, ax=ax)
```





```
[20]: def remove_outliers(df, columns):
    for col in columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df[(df[col] < lower_bound)] = pd.NA
        df[(df[col] > upper_bound)] = pd.NA
    return df

[21]: diabetes_x_resampled = remove_outliers(diabetes_x_resampled,
    ↪diabetes_x_resampled.columns)

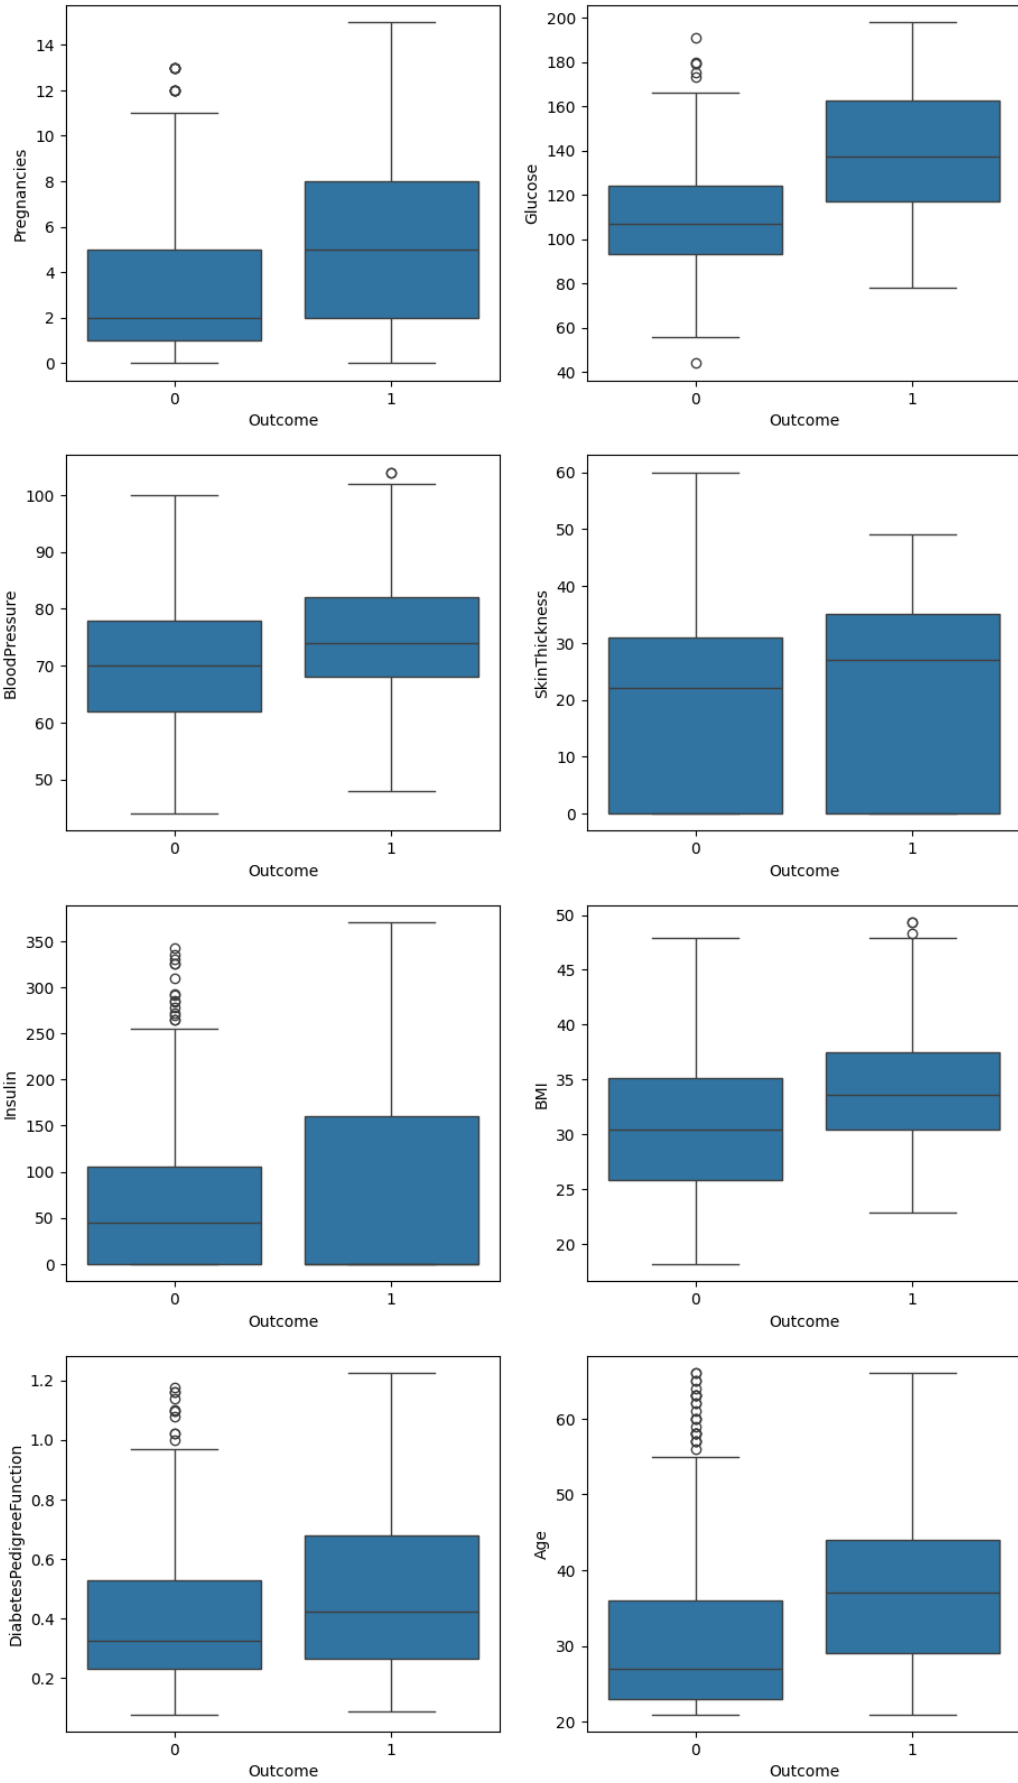
[22]: diabetes_x_copy = remove_outliers(diabetes_x_copy, diabetes_x_copy.columns)

[23]: combined_df = pd.concat([diabetes_x_resampled, diabetes_y_resampled],
    ↪axis=1)
combined_df_cp = pd.concat([diabetes_x_copy, diabetes_y_copy], axis=1)

[24]: combined_df.dropna(inplace=True)
combined_df_cp.dropna(inplace=True)

[25]: diabetes_x = combined_df.drop("Outcome", axis=1)
diabetes_y = combined_df["Outcome"]
diabetes_x_copy = combined_df_cp.drop("Outcome", axis=1)
diabetes_y_copy = combined_df_cp["Outcome"]

[33]: fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(11, 20))
for i, col in enumerate(diabetes_x.columns):
    ax = axes[i // 2, i % 2]
    sns.boxplot(data=diabetes_x, x=diabetes_y, y=col, ax=ax)
```




```
[34]: print(f"After removing outliers, there are {diabetes_x.shape[0]} samples_
↳left.")
print(f"There are {diabetes_x[diabetes_y == 1].shape[0]} positive cases_
↳and {diabetes_x[diabetes_y== 0].shape[0]} negative cases")
```

After removing outliers, there are 835 samples left.  
There are 394 positive cases and 441 negative cases

Correlation matrix of balanced dataset

```
[43]: corr = combined_df.corr()
plt.figure(figsize=(9, 9))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm", square=True)
```


[43]: <Axes: >



Correlation matrix of unbalanced dataset

```
[45]: corr_cp = combined_df_cp.corr()  
plt.figure(figsize=(9, 9))  
sns.heatmap(corr_cp, annot=True, fmt=".2f", cmap="coolwarm", square=True)
```

[45]: <Axes: >



```
[46]: scaler = StandardScaler()  
diabetes_X_scaled = scaler.fit_transform(diabetes_x)  
diabetes_X_scaled_cp = scaler.fit_transform(diabetes_x_copy)
```

```
[47]: X_train, X_test, y_train, y_test = train_test_split(diabetes_X_scaled, ↵  
    ↪diabetes_y, test_size=0.2, random_state=42)
```

```
[48]: X_train_02, X_test_02, y_train_02, y_test_02 = ↵  
    ↪train_test_split(diabetes_X_scaled_cp, diabetes_y_copy, test_size=0.2, ↵  
    ↪random_state=42)
```

Training a Random Forest Classifier on the balanced dataset

```
[50]: rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
rf_model.fit(X_train, y_train)
```

```
[50]: RandomForestClassifier(random_state=42)
```

```
[51]: y_pred = rf_model.predict(X_test)
```

```
[52]: accuracy = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy)
```


Accuracy: 0.8802395209580839

```
[53]: confusion_mat = confusion_matrix(y_test, y_pred)  
print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[75 13]  
 [ 7 72]]
```

```
[56]: plt.figure(figsize=(6, 6))  
sns.heatmap(confusion_mat, annot=True, annot_kws={"size": 16}, fmt='d', ↵  
    ↪cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No ↵  
    ↪Diabetes", "Diabetes"], square=True)  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.title("Confusion Matrix (Random Forest)")  
plt.show()
```



```
[57]: print("Classification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.85	0.88	88
1	0.85	0.91	0.88	79
accuracy			0.88	167
macro avg	0.88	0.88	0.88	167
weighted avg	0.88	0.88	0.88	167

Training a Random Forest Classifier on the unbalanced dataset

```
[59]: rf_model_02 = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model_02.fit(X_train_02, y_train_02)
```

```
[59]: RandomForestClassifier(random_state=42)
```

```
[60]: y_pred_02 = rf_model_02.predict(X_test_02)
```

```
[61]: accuracy_02 = accuracy_score(y_test_02, y_pred_02)
print("Accuracy on unbalanced dataset:", accuracy_02)
```

Accuracy on unbalanced dataset: 0.7421875


```
[62]: y_test_02
```

```
[62]: 338    1
      763    0
      104    0
      437    0
      184    0
      ..
      723    0
      246    0
      346    0
      275    0
      358    0
      Name: Outcome, Length: 128, dtype: int64
```

```
[63]: confusion_matrix_02 = confusion_matrix(y_test_02, y_pred_02)
print("Confusion Matrix (Unbalanced Dataset):")
print(confusion_matrix_02)
```

Confusion Matrix (Unbalanced Dataset):  
[[79 13]  
 [20 16]]

```
[64]: plt.figure(figsize=(6,6))
sns.heatmap(confusion_matrix_02, annot=True, annot_kws={"size": 16},
            fmt='d', cmap='Blues', xticklabels=["No Diabetes", "Diabetes"],
            yticklabels=["No Diabetes", "Diabetes"], square=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix (Random Forest)")
plt.show()
```



```
[65]: print("Classification Report:\n", classification_report(y_test_02,
→y_pred_02))
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.86	0.83	92
1	0.55	0.44	0.49	36
accuracy			0.74	128
macro avg	0.67	0.65	0.66	128
weighted avg	0.73	0.74	0.73	128

Training a Multi-Layer Perceptron on the balanced dataset

```
[142]: class DiabetesModel(nn.Module):
        def __init__(self, input_size=8, hidden_size=14, output_size=1):
```



```

    super(DiabetesModel, self).__init__()
    self.fc1 = nn.Linear(input_size, hidden_size)
    self.fc2 = nn.Linear(hidden_size, hidden_size)
    self.out = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        x = F.leaky_relu(self.fc1(x))
        x = F.leaky_relu(self.fc2(x))
        x = self.out(x)
        return x

```

```

[143]: torch.manual_seed(42)
nn_model_01 = DiabetesModel()
nn_model_01
nn_model_01.to(device=device)

```

```

[143]: DiabetesModel(
  (fc1): Linear(in_features=8, out_features=14, bias=True)
  (fc2): Linear(in_features=14, out_features=14, bias=True)
  (out): Linear(in_features=14, out_features=1, bias=True)
)

```

```

[144]: torch.manual_seed(42)
nn_model_02 = DiabetesModel()
nn_model_02
nn_model_02.to(device=device)

```

```

[144]: DiabetesModel(
  (fc1): Linear(in_features=8, out_features=14, bias=True)
  (fc2): Linear(in_features=14, out_features=14, bias=True)
  (out): Linear(in_features=14, out_features=1, bias=True)
)

```

We are using cross entropy loss for calculating the loss of the model, and NADAM for optimizing the weights and biases of the model

```

[145]: loss_fn = nn.BCEWithLogitsLoss()
optimizer = NAdam(nn_model_01.parameters(), lr=0.01)

```

```

[146]: loss_fn2 = nn.BCEWithLogitsLoss()
optimizer2 = NAdam(nn_model_02.parameters(), lr=0.01)

```

```

[147]: X_train_tensor = torch.tensor(X_train, dtype=torch.float32, device=device)
y_train_tensor = torch.tensor(y_train.values, dtype=torch.float32, ↵
↪device=device).view(-1, 1)

```

```
X_test_tensor = torch.tensor(X_test, dtype=torch.float32, device=device)
y_test_tensor = torch.tensor(y_test.values, dtype=torch.float32,
    ↪device=device).view(-1, 1)
```

```
[148]: X_train_tensor1 = torch.tensor(X_train_02, dtype=torch.float32,
    ↪device=device)
y_train_tensor1 = torch.tensor(y_train_02.values, dtype=torch.float32,
    ↪device=device).view(-1, 1)
X_test_tensor1 = torch.tensor(X_test_02, dtype=torch.float32,
    ↪device=device)
y_test_tensor1 = torch.tensor(y_test_02.values, dtype=torch.float32,
    ↪device=device).view(-1, 1)
```

Training the neural network on the balanced dataset

```
[149]: epochs = 1000
best_model_loss = float('inf')
best_model_weights = None
patience = 7

loss_list = []
accuracy_list = []

for i in range(epochs):
    # train_loss = 0.0
    nn_model_01.train()

    outputs = nn_model_01(X_train_tensor)
    loss = loss_fn(outputs, y_train_tensor)
    loss_list.append(loss.item())

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    nn_model_01.eval()
    with torch.no_grad():
        outputs = nn_model_01(X_test_tensor)
        test_loss = loss_fn(outputs, y_test_tensor)
    y_pred = torch.sigmoid(outputs).round()
    accuracy = (y_pred == y_test_tensor).float().mean().item()
    accuracy_list.append(accuracy)

    if test_loss < best_model_loss:
        best_model_loss = test_loss
```

```

        best_model_weights = copy.deepcopy(nn_model_01.state_dict())
        patience_counter = 0
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print(f"Early stopping at epoch {i+1}")
            break
        print(f"Epoch {i+1}/{epochs}, Train Loss: {loss.item():.4f}, Val Loss:  
         {test_loss.item():.4f}")

```

```

Epoch 1/1000, Train Loss: 0.6923, Val Loss: 0.6869
Epoch 2/1000, Train Loss: 0.6869, Val Loss: 0.6823
Epoch 3/1000, Train Loss: 0.6820, Val Loss: 0.6764
Epoch 4/1000, Train Loss: 0.6764, Val Loss: 0.6688
Epoch 5/1000, Train Loss: 0.6694, Val Loss: 0.6588
Epoch 6/1000, Train Loss: 0.6604, Val Loss: 0.6463
Epoch 7/1000, Train Loss: 0.6490, Val Loss: 0.6315
Epoch 8/1000, Train Loss: 0.6355, Val Loss: 0.6143
Epoch 9/1000, Train Loss: 0.6199, Val Loss: 0.5949
Epoch 10/1000, Train Loss: 0.6030, Val Loss: 0.5750
Epoch 11/1000, Train Loss: 0.5853, Val Loss: 0.5558
Epoch 12/1000, Train Loss: 0.5682, Val Loss: 0.5389
Epoch 13/1000, Train Loss: 0.5522, Val Loss: 0.5242
Epoch 14/1000, Train Loss: 0.5377, Val Loss: 0.5100
Epoch 15/1000, Train Loss: 0.5244, Val Loss: 0.4983
Epoch 16/1000, Train Loss: 0.5125, Val Loss: 0.4879
Epoch 17/1000, Train Loss: 0.5018, Val Loss: 0.4791
Epoch 18/1000, Train Loss: 0.4919, Val Loss: 0.4711
Epoch 19/1000, Train Loss: 0.4831, Val Loss: 0.4650
Epoch 20/1000, Train Loss: 0.4754, Val Loss: 0.4592
Epoch 21/1000, Train Loss: 0.4683, Val Loss: 0.4551
Epoch 22/1000, Train Loss: 0.4621, Val Loss: 0.4507
Epoch 23/1000, Train Loss: 0.4567, Val Loss: 0.4507
Epoch 24/1000, Train Loss: 0.4522, Val Loss: 0.4470
Epoch 25/1000, Train Loss: 0.4484, Val Loss: 0.4493
Epoch 26/1000, Train Loss: 0.4449, Val Loss: 0.4447
Epoch 27/1000, Train Loss: 0.4420, Val Loss: 0.4546
Epoch 28/1000, Train Loss: 0.4407, Val Loss: 0.4447
Epoch 29/1000, Train Loss: 0.4410, Val Loss: 0.4588
Epoch 30/1000, Train Loss: 0.4378, Val Loss: 0.4443
Epoch 31/1000, Train Loss: 0.4347, Val Loss: 0.4543
Epoch 32/1000, Train Loss: 0.4310, Val Loss: 0.4452
Epoch 33/1000, Train Loss: 0.4284, Val Loss: 0.4515
Epoch 34/1000, Train Loss: 0.4260, Val Loss: 0.4448
Epoch 35/1000, Train Loss: 0.4240, Val Loss: 0.4508

```

Epoch 36/1000, Train Loss: 0.4219, Val Loss: 0.4430  
Epoch 37/1000, Train Loss: 0.4201, Val Loss: 0.4496  
Epoch 38/1000, Train Loss: 0.4185, Val Loss: 0.4418  
Epoch 39/1000, Train Loss: 0.4170, Val Loss: 0.4500  
Epoch 40/1000, Train Loss: 0.4154, Val Loss: 0.4406  
Epoch 41/1000, Train Loss: 0.4141, Val Loss: 0.4509  
Epoch 42/1000, Train Loss: 0.4127, Val Loss: 0.4395  
Epoch 43/1000, Train Loss: 0.4117, Val Loss: 0.4519  
Epoch 44/1000, Train Loss: 0.4096, Val Loss: 0.4389  
Epoch 45/1000, Train Loss: 0.4085, Val Loss: 0.4501  
Epoch 46/1000, Train Loss: 0.4065, Val Loss: 0.4391  
Epoch 47/1000, Train Loss: 0.4052, Val Loss: 0.4487  
Epoch 48/1000, Train Loss: 0.4036, Val Loss: 0.4384  
Epoch 49/1000, Train Loss: 0.4028, Val Loss: 0.4481  
Epoch 50/1000, Train Loss: 0.4012, Val Loss: 0.4380  
Epoch 51/1000, Train Loss: 0.4000, Val Loss: 0.4475  
Epoch 52/1000, Train Loss: 0.3981, Val Loss: 0.4370  
Epoch 53/1000, Train Loss: 0.3964, Val Loss: 0.4458  
Epoch 54/1000, Train Loss: 0.3949, Val Loss: 0.4355  
Epoch 55/1000, Train Loss: 0.3940, Val Loss: 0.4448  
Epoch 56/1000, Train Loss: 0.3924, Val Loss: 0.4342  
Epoch 57/1000, Train Loss: 0.3914, Val Loss: 0.4448  
Epoch 58/1000, Train Loss: 0.3905, Val Loss: 0.4322  
Epoch 59/1000, Train Loss: 0.3901, Val Loss: 0.4444  
Epoch 60/1000, Train Loss: 0.3891, Val Loss: 0.4324  
Epoch 61/1000, Train Loss: 0.3894, Val Loss: 0.4460  
Epoch 62/1000, Train Loss: 0.3875, Val Loss: 0.4316  
Epoch 63/1000, Train Loss: 0.3868, Val Loss: 0.4446  
Epoch 64/1000, Train Loss: 0.3850, Val Loss: 0.4306  
Epoch 65/1000, Train Loss: 0.3848, Val Loss: 0.4442  
Epoch 66/1000, Train Loss: 0.3825, Val Loss: 0.4299  
Epoch 67/1000, Train Loss: 0.3805, Val Loss: 0.4418  
Epoch 68/1000, Train Loss: 0.3794, Val Loss: 0.4301  
Epoch 69/1000, Train Loss: 0.3789, Val Loss: 0.4418  
Epoch 70/1000, Train Loss: 0.3769, Val Loss: 0.4303  
Epoch 71/1000, Train Loss: 0.3762, Val Loss: 0.4434  
Epoch 72/1000, Train Loss: 0.3755, Val Loss: 0.4307  
Early stopping at epoch 73

```
[150]: nn_model_01.to(device="cpu")
nn_model_01.load_state_dict(best_model_weights)
nn_model_01.eval()

X_test_tensor = X_test_tensor.to(device="cpu")
```

```
with torch.no_grad():
    outputs = nn_model_01(X_test_tensor)
    y_pred = torch.sigmoid(outputs).round().numpy()
```

```
[151]: y_test_tensor = y_test_tensor.to(device="cpu")

print("Classification Report:")
print(classification_report(y_test_tensor, y_pred))
```


```
Classification Report:
              precision    recall  f1-score   support

     0.0         0.86      0.84      0.85         88
     1.0         0.83      0.85      0.84         79

 accuracy              0.84         167
 macro avg           0.84      0.84      0.84         167
weighted avg           0.84      0.84      0.84         167
```

```
[152]: conf_matrix = confusion_matrix(y_test_tensor, y_pred)
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=['No Diabetes', 'Diabetes'], yticklabels=['No Diabetes', 'Diabetes'], square=True)
```


```
[152]: <Axes: >
```



```
[153]: print(f"Accuracy: {accuracy_score(y_test_tensor, y_pred):.4f}")
```

Accuracy: 0.8443

```
[156]: plt.figure(figsize=(10, 5))
sns.lineplot(x=range(len(loss_list)), y=loss_list, label='Train Loss')
sns.lineplot(x=range(len(accuracy_list)), y=accuracy_list, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss / Accuracy')
plt.title('Loss and Accuracy over Epochs')
plt.legend()
plt.show()
```



Training the neural network on the unbalanced dataset

```
[157]: epochs = 400
best_model_loss1 = float('inf')
best_model_weights1 = None
patience = 7

loss_list1 = []
accuracy_list1 = []
```

```

for i in range(epochs):
    # train_loss = 0.0
    nn_model_02.train()

    outputs1 = nn_model_02(X_train_tensor1)
    loss1 = loss_fn2(outputs1, y_train_tensor1)
    loss_list1.append(loss1.item())

    optimizer2.zero_grad()
    loss1.backward()
    optimizer2.step()

    nn_model_02.eval()
    with torch.no_grad():
        outputs1 = nn_model_02(X_test_tensor1)
        test_loss1 = loss_fn2(outputs1, y_test_tensor1)
    y_pred = torch.sigmoid(outputs1).round()
    accuracy = (y_pred == y_test_tensor1).float().mean().item()
    accuracy_list1.append(accuracy)

    if test_loss1 < best_model_loss1:
        best_model_loss1 = test_loss1
        best_model_weights1 = copy.deepcopy(nn_model_02.state_dict())
        patience_counter1 = 0
    else:
        patience_counter1 += 1
        if patience_counter1 >= patience:
            print(f"Early stopping at epoch {i+1}")
            break
    print(f"Epoch {i+1}/{epochs}, Train Loss: {loss1.item():.4f}, Val Loss:
→ {test_loss1.item():.4f}")

```

```

Epoch 1/400, Train Loss: 0.6982, Val Loss: 0.6940
Epoch 2/400, Train Loss: 0.6881, Val Loss: 0.6849
Epoch 3/400, Train Loss: 0.6794, Val Loss: 0.6753
Epoch 4/400, Train Loss: 0.6702, Val Loss: 0.6646
Epoch 5/400, Train Loss: 0.6596, Val Loss: 0.6520
Epoch 6/400, Train Loss: 0.6469, Val Loss: 0.6371
Epoch 7/400, Train Loss: 0.6317, Val Loss: 0.6201
Epoch 8/400, Train Loss: 0.6138, Val Loss: 0.6022
Epoch 9/400, Train Loss: 0.5938, Val Loss: 0.5833
Epoch 10/400, Train Loss: 0.5725, Val Loss: 0.5659
Epoch 11/400, Train Loss: 0.5512, Val Loss: 0.5519
Epoch 12/400, Train Loss: 0.5309, Val Loss: 0.5404

```

```
Epoch 13/400, Train Loss: 0.5126, Val Loss: 0.5316
Epoch 14/400, Train Loss: 0.4969, Val Loss: 0.5252
Epoch 15/400, Train Loss: 0.4829, Val Loss: 0.5201
Epoch 16/400, Train Loss: 0.4703, Val Loss: 0.5167
Epoch 17/400, Train Loss: 0.4592, Val Loss: 0.5143
Epoch 18/400, Train Loss: 0.4492, Val Loss: 0.5122
Epoch 19/400, Train Loss: 0.4401, Val Loss: 0.5112
Epoch 20/400, Train Loss: 0.4317, Val Loss: 0.5104
Epoch 21/400, Train Loss: 0.4238, Val Loss: 0.5094
Epoch 22/400, Train Loss: 0.4161, Val Loss: 0.5070
Epoch 23/400, Train Loss: 0.4087, Val Loss: 0.5066
Epoch 24/400, Train Loss: 0.4025, Val Loss: 0.5068
Epoch 25/400, Train Loss: 0.3970, Val Loss: 0.5074
Epoch 26/400, Train Loss: 0.3921, Val Loss: 0.5081
Epoch 27/400, Train Loss: 0.3879, Val Loss: 0.5087
Epoch 28/400, Train Loss: 0.3842, Val Loss: 0.5107
Epoch 29/400, Train Loss: 0.3809, Val Loss: 0.5101
Early stopping at epoch 30
```

```
[158]: nn_model_02.load_state_dict(best_model_weights1)
nn_model_02.to(device="cpu")
X_test_tensor1 = X_test_tensor1.to(device="cpu")

nn_model_02.eval()
with torch.no_grad():
    outputs1 = nn_model_02(X_test_tensor1)
    y_pred_02 = torch.sigmoid(outputs1).round().numpy()
```

```
[159]: y_test_tensor1 = y_test_tensor1.to(device="cpu")

print("Classification Report:")
print(classification_report(y_test_tensor1, y_pred_02))
```

Classification Report:


	precision	recall	f1-score	support
0.0	0.80	0.88	0.84	92
1.0	0.59	0.44	0.51	36
accuracy			0.76	128
macro avg	0.70	0.66	0.67	128
weighted avg	0.74	0.76	0.75	128

```
[160]: print(f"Accuracy: {accuracy_score(y_test_tensor1, y_pred_02):.4f}")
```



Accuracy: 0.7578

```
[161]: plt.figure(figsize=(10, 5))
sns.lineplot(x=range(len(loss_list1)), y=loss_list1, label='Train Loss')
sns.lineplot(x=range(len(accuracy_list1)), y=accuracy_list1, label='Test_
→Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss / Accuracy')
plt.title('Loss and Accuracy over Epochs')
plt.legend()
plt.show()
```



Training a Support Vector Machine on the balanced dataset

```
[162]: svc_model = SVC(random_state=42, kernel='rbf')
```

```
[163]: svc_model.fit(X_train, y_train)
```

```
[163]: SVC(random_state=42)
```

```
[164]: y_pred = svc_model.predict(X_test)
```

```
[165]: accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
```

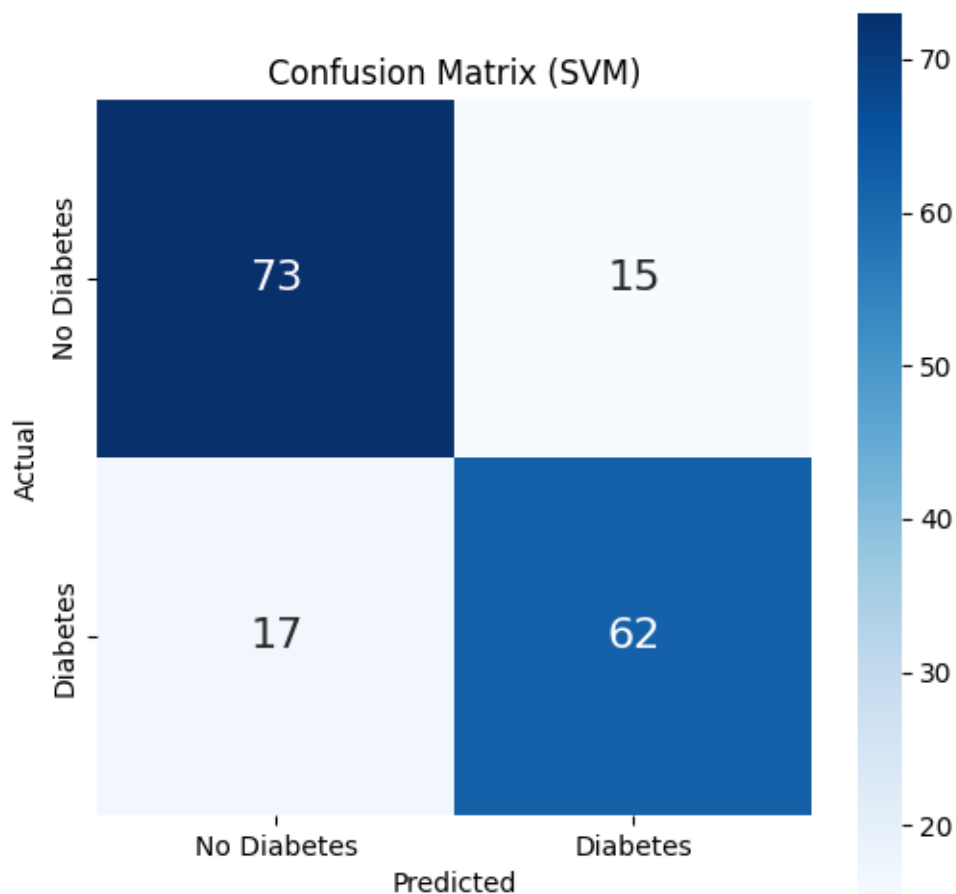
Accuracy: 0.8084

```
[166]: confusion_mat = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[73 15]
 [17 62]]
```

```
[167]: plt.figure(figsize=(6, 6))
sns.heatmap(confusion_mat, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No Diabetes", "Diabetes"], square=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix (SVM)")
plt.show()
```



```
[168]: print("Classification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.83	0.82	88
1	0.81	0.78	0.79	79
accuracy			0.81	167
macro avg	0.81	0.81	0.81	167
weighted avg	0.81	0.81	0.81	167

Training a Support Vector Machine on the unbalanced dataset.

```
[169]: svc_model_02 = SVC(random_state=42, kernel='rbf')
```

```
[170]: svc_model_02.fit(X_train_02, y_train_02)
```

```
[170]: SVC(random_state=42)
```

```
[171]: y_pred_02 = svc_model_02.predict(X_test_02)
```

```
[172]: accuracy_02 = accuracy_score(y_test_02, y_pred_02)
print(f"Accuracy: {accuracy_02:.4f}")
```


Accuracy: 0.7578

```
[173]: confusion_mat = confusion_matrix(y_test_02, y_pred_02)
print("Confusion Matrix:\n", confusion_mat)
```

Confusion Matrix:

```
[[81 11]
 [20 16]]
```

```
[174]: plt.figure(figsize=(6, 6))
sns.heatmap(confusion_mat, annot=True, annot_kws={"size": 16}, fmt='d',
            cmap='Blues', xticklabels=["No Diabetes", "Diabetes"], yticklabels=["No
            Diabetes", "Diabetes"], square=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix (SVM)")
plt.show()
```



```
[175]: print("Classification Report:\n", classification_report(y_test_02,
→y_pred_02))
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.88	0.84	92
1	0.59	0.44	0.51	36
accuracy			0.76	128
macro avg	0.70	0.66	0.67	128
weighted avg	0.74	0.76	0.75	128

Plotting the ROC-AUC curves of the models trained on the balanced dataset

```
[176]: y_pred_rf = rf_model.predict_proba(X_test)[:, 1]
y_pred_svc = svc_model.decision_function(X_test)
```


```
y_pred_nn = torch.sigmoid(nn_model_01(X_test_tensor)).detach().numpy()
```

```
[179]: plt.figure(figsize=(7, 7))

fpr1, tpr1, _ = roc_curve(y_test, y_pred_rf)
fpr2, tpr2, _ = roc_curve(y_test, y_pred_svc)
fpr3, tpr3, _ = roc_curve(y_test, y_pred_nn)

roc_auc_rf = auc(fpr1, tpr1)
roc_auc_svc = auc(fpr2, tpr2)
roc_auc_nn = auc(fpr3, tpr3)

plt.plot(fpr1, tpr1, color='blue', label=f'Random Forest (AUC = {roc_auc_rf:.2f})')
plt.plot(fpr2, tpr2, color='green', label=f'SVM (AUC = {roc_auc_svc:.2f})')
plt.plot(fpr3, tpr3, color='red', label=f'Neural Network (AUC = {roc_auc_nn:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



```
[180]: nn_model_01.eval()
with torch.no_grad():
    outputs = nn_model_01(X_test_tensor)
    y_pred = torch.sigmoid(outputs).round().numpy()

results_df = pd.DataFrame({
    'Model': ['Random Forest', 'Neural Network', 'SVM'],
    'Accuracy': [
        accuracy_score(y_test, rf_model.predict(X_test)),
        accuracy_score(y_test, y_pred),
        accuracy_score(y_test, svc_model.predict(X_test)),
    ],
})
```

```

'AUC': [roc_auc_rf, roc_auc_svc, roc_auc_nn],
'F1-Score': [
    f1_score(y_test, rf_model.predict(X_test)),
    f1_score(y_test, y_pred),
    f1_score(y_test, svc_model.predict(X_test)),
],
'Precision': [
    precision_score(y_test, rf_model.predict(X_test)),
    precision_score(y_test, y_pred),
    precision_score(y_test, svc_model.predict(X_test)),
],
'Recall': [
    recall_score(y_test, rf_model.predict(X_test)),
    recall_score(y_test, y_pred),
    recall_score(y_test, svc_model.predict(X_test)),
]
})

```

[181]: results\_df

	Model	Accuracy	AUC	F1-Score	Precision	Recall
0	Random Forest	0.880240	0.946778	0.878049	0.847059	0.911392
1	Neural Network	0.844311	0.889097	0.837500	0.827160	0.848101
2	SVM	0.808383	0.886939	0.794872	0.805195	0.784810